

EE559 HW6

July 15, 2018

1 1. Multi-class and Multi-Label Classification Using Support Vector Machines

2 Name:Yuxiang Hou

3 USCID:1126739072

Each instance has three labels: Families, Genus, and Species. Each of the labels has multiple classes. We wish to solve a multi-class and multi-label problem. One of the most important approaches to multi-class classification is to train a classifier for each label. We first try this approach:

- i. Research exact match and hamming score/ loss methods for evaluating multi- label classification and use them in evaluating the classifiers in this problem.

1.The hamming scores means that the percent of misclassification.

- ii. Train a SVM for each of the labels, using Gaussian kernels and one versus all classifiers. Determine the weight of the SVM penalty and the width of the Gaussian Kernel using 10 fold cross validation.¹ You are welcome to try to solve the problem with both normalized and raw attributes and report the results.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

In [2]: data=pd.read_csv("Frogs_MFCCs.csv")

In [3]: from sklearn import preprocessing
from sklearn import svm

In [4]: def trasfer_data():
    le = preprocessing.LabelEncoder()
    le.fit(data["Family"])
    data["Family"]=le.transform(data["Family"])

    le.fit(data["Genus"])
    data["Genus"]=le.transform(data["Genus"])
```

```
le.fit(data["Species"])
data["Species"]=le.transform(data["Species"])
```

```
In [5]: trasfer_data()
```

```
In [6]: linear_svc=svm.SVC(kernel='rbf')
```

```
In [7]: family_data=data.drop(["Genus","Species","RecordID"],axis=1)
        genus_data=data.drop(["Family","Species","RecordID"],axis=1)
        species_data=data.drop(["Family","Genus","RecordID"],axis=1)
```

```
In [8]: family_y=list(data["Family"])
        family_x=data.drop(["Family","Genus","Species","RecordID"],axis=1)

        genus_y=list(data["Genus"])
        genus_x=data.drop(["Family","Genus","Species","RecordID"],axis=1)

        species_y=list(data["Species"])
        species_x=data.drop(["Family","Genus","Species","RecordID"],axis=1)
```

```
In [9]: from sklearn.metrics import accuracy_score
```

Choose 70% of the data randomly as the training set.

```
In [10]: from sklearn.model_selection import train_test_split
```

```
In [11]: family_x_train, family_x_test,family_y_train,family_y_test = train_test_split(family_x, family_y,
        genus_x_train, genus_x_test, genus_y_train, genus_y_test = train_test_split(genus_x, genus_y,
        species_x_train, species_x_test, species_y_train, species_y_test = train_test_split(species_x, species_y,
```

4 Raw SVM Gaussian Kernels

```
In [50]: from sklearn.model_selection import GridSearchCV
```

```
In [51]: gamma_range = np.logspace(-3, 1, 5)
```

```
In [52]: linear_svc=svm.SVC(kernel='rbf') #there are two main parameters, C and degree
        parameters = { 'C':[1,2,3,4], 'gamma':[x for x in gamma_range]}
        clf = GridSearchCV(linear_svc, parameters,cv=10)
        clf.fit(family_x_train,family_y_train)
```

```
Out[52]: GridSearchCV(cv=10, error_score='raise',
        estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
        decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
        max_iter=-1, probability=False, random_state=None, shrinking=True,
        tol=0.001, verbose=False),
        fit_params=None, iid=True, n_jobs=1,
        param_grid={'C': [1, 2, 3, 4], 'gamma': [0.001, 0.01, 0.1, 1.0, 10.0]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
        scoring=None, verbose=0)
```

```

In [53]: res=clf.predict(family_x_test)
In [54]: accuracy_score(res,family_y_test)
Out[54]: 0.9884205650764243
In [55]: clf.fit(genus_x_train,genus_y_train)
          genus_res=clf.predict(genus_x_train)
          print(accuracy_score(genus_res,genus_y_train))
0.9960285941223193

In [56]: clf.fit(species_x_train,species_y_train)
          species_res=clf.predict(species_x_train)
          print(accuracy_score(species_res,species_y_train))
0.9978157267672756

```

From the above result, we can find that if we use the Raw data without the Normalization, the accuracy score is around 99% percent, and then we will try to utilize the normalized result.

```

In [ ]: from sklearn.preprocessing import MinMaxScaler
In [57]: min_max_scaler = preprocessing.MinMaxScaler()
          family_x_train_normalized=min_max_scaler.fit_transform(family_x_train)
          genus_x_train_normalized=min_max_scaler.fit_transform(genus_x_train)
          species_x_train_normalized=min_max_scaler.fit_transform(species_x_train)
In [58]: clf.fit(family_x_train_normalized,family_y_train)
          family_res=clf.predict(family_x_train_normalized)
          print(accuracy_score(family_res,family_y_train))
1.0

```

After normalization, the accuracy improved from 98.8% to 100% percent for family class

```

In [59]: clf.fit(genus_x_train_normalized,genus_y_train)
          genus_res=clf.predict(genus_x_train_normalized)
          print(accuracy_score(genus_res,genus_y_train))
0.9944400317712471

```

After normalization, the accuracy of genus class decreased from 99.6% to 99.4% percent.

```

In [60]: clf.fit(species_x_train_normalized,species_y_train)
          species_res=clf.predict(species_x_train_normalized)
          print(accuracy_score(species_res,species_y_train))
0.9950357426528992

```

Conclusion: From the above experiments, we can find that the normalized result will improve the results from 99.7% percent to 99.5% percent

5 iii. Repeat 1(b)ii with L1-penalized SVMs. Remember to normalize the at-tributes.

```
In [80]: svm_model=svm.LinearSVC(penalty="l1")
```

```
In [88]: # parameters = { 'C':[1,2,3,4,5,6], 'tol':[0.0001,0.001,0.01]}  
#then we have to utilize the for loop to get it
```

```
def L1_SVM_Model(x,y):  
    best_c=None  
    best_tol=None  
    res=0  
    for i in range(1,11):  
        C=i  
        for j in [0.0001,0.001, 0.01,0.1]:  
            tol=j  
            svm_model=svm.LinearSVC(penalty="l1",C=C,tol=tol,dual=False)  
            scores = cross_val_score(svm_model,x,y,cv=10)  
            avg_scores=np.mean(scores)  
            print("When C = "+str(C)+" And tol= "+str(tol)+" The result is: "+str(avg_scores))  
  
            if avg_scores>res:  
                res=avg_scores  
                best_c=C  
                best_tol=tol  
    return res,best_c,best_tol
```

```
In [89]: print(L1_SVM_Model(family_x_train_normalized,family_y_train))  
print(L1_SVM_Model(genus_x_train_normalized,genus_y_train))  
print(L1_SVM_Model(species_x_train_normalized,species_y_train))
```

```
When C = 1 And tol= 0.0001 The result is: 0.936466680372642  
When C = 1 And tol= 0.001 The result is: 0.9378555739481682  
When C = 1 And tol= 0.01 The result is: 0.9354730390459987  
When C = 1 And tol= 0.1 The result is: 0.928521475627182  
When C = 2 And tol= 0.0001 The result is: 0.9368627199766024  
When C = 2 And tol= 0.001 The result is: 0.93706152713366  
When C = 2 And tol= 0.01 The result is: 0.9356710588479789  
When C = 2 And tol= 0.1 The result is: 0.9285242259207338  
When C = 3 And tol= 0.0001 The result is: 0.93706152713366  
When C = 3 And tol= 0.001 The result is: 0.93706152713366  
When C = 3 And tol= 0.01 The result is: 0.9366650946332671  
When C = 3 And tol= 0.1 The result is: 0.9277266414575962  
When C = 4 And tol= 0.0001 The result is: 0.9374591414477752  
When C = 4 And tol= 0.001 The result is: 0.93706152713366  
When C = 4 And tol= 0.01 The result is: 0.9362682676742292  
When C = 4 And tol= 0.1 The result is: 0.9269329859865703
```

When C = 5 And tol= 0.0001 The result is: 0.93706152713366
 When C = 5 And tol= 0.001 The result is: 0.9374591414477752
 When C = 5 And tol= 0.01 The result is: 0.9366643072781896
 When C = 5 And tol= 0.1 The result is: 0.9279266273132134
 When C = 6 And tol= 0.0001 The result is: 0.9372603342907176
 When C = 6 And tol= 0.001 The result is: 0.9372603342907176
 When C = 6 And tol= 0.01 The result is: 0.9360686731620941
 When C = 6 And tol= 0.1 The result is: 0.9277286043960705
 When C = 7 And tol= 0.0001 The result is: 0.9374591414477752
 When C = 7 And tol= 0.001 The result is: 0.9372603342907176
 When C = 7 And tol= 0.01 The result is: 0.9364670748312868
 When C = 7 And tol= 0.1 The result is: 0.9273325648013723
 When C = 8 And tol= 0.0001 The result is: 0.9372611216457951
 When C = 8 And tol= 0.001 The result is: 0.9376579486048329
 When C = 8 And tol= 0.01 The result is: 0.9364662874762095
 When C = 8 And tol= 0.1 The result is: 0.9277266399046461
 When C = 9 And tol= 0.0001 The result is: 0.9374591414477752
 When C = 9 And tol= 0.001 The result is: 0.9374591414477752
 When C = 9 And tol= 0.01 The result is: 0.9362682676742292
 When C = 9 And tol= 0.1 The result is: 0.9275286185404532
 When C = 10 And tol= 0.0001 The result is: 0.9372611216457951
 When C = 10 And tol= 0.001 The result is: 0.9372611216457951
 When C = 10 And tol= 0.01 The result is: 0.9364662874762095
 When C = 10 And tol= 0.1 The result is: 0.9273306003006855
 (0.9378555739481682, 1, 0.001)
 When C = 1 And tol= 0.0001 The result is: 0.939441964582055
 When C = 1 And tol= 0.001 The result is: 0.9390455227360268
 When C = 1 And tol= 0.01 The result is: 0.9384498854674241
 When C = 1 And tol= 0.1 The result is: 0.9317128733315044
 When C = 2 And tol= 0.0001 The result is: 0.942616990354242
 When C = 2 And tol= 0.001 The result is: 0.9420201681288202
 When C = 2 And tol= 0.01 The result is: 0.9412217837383643
 When C = 2 And tol= 0.1 The result is: 0.9328974558803198
 When C = 3 And tol= 0.0001 The result is: 0.9447995612444234
 When C = 3 And tol= 0.001 The result is: 0.9444023335404728
 When C = 3 And tol= 0.01 The result is: 0.9426177761283606
 When C = 3 And tol= 0.1 The result is: 0.9344910564940039
 When C = 4 And tol= 0.0001 The result is: 0.9447987707555854
 When C = 4 And tol= 0.001 The result is: 0.9451971771300874
 When C = 4 And tol= 0.01 The result is: 0.9432106661443276
 When C = 4 And tol= 0.1 The result is: 0.933696241034435
 When C = 5 And tol= 0.0001 The result is: 0.9455932229459221
 When C = 5 And tol= 0.001 The result is: 0.9449971865877588
 When C = 5 And tol= 0.01 The result is: 0.9436098552429915
 When C = 5 And tol= 0.1 The result is: 0.9360744555641087
 When C = 6 And tol= 0.0001 The result is: 0.9459924104823736
 When C = 6 And tol= 0.001 The result is: 0.9461900373879214
 When C = 6 And tol= 0.01 The result is: 0.9449948244668793

When C = 6 And tol= 0.1 The result is: 0.9350914258252814
 When C = 7 And tol= 0.0001 The result is: 0.9463896381863245
 When C = 7 And tol= 0.001 The result is: 0.9467856730755653
 When C = 7 And tol= 0.01 The result is: 0.9445995800664768
 When C = 7 And tol= 0.1 The result is: 0.9372688764153537
 When C = 8 And tol= 0.0001 The result is: 0.9471844770612193
 When C = 8 And tol= 0.001 The result is: 0.9467880445608264
 When C = 8 And tol= 0.01 The result is: 0.9453932449017362
 When C = 8 And tol= 0.1 The result is: 0.934283617844524
 When C = 9 And tol= 0.0001 The result is: 0.9457955819141448
 When C = 9 And tol= 0.001 The result is: 0.9465896318624137
 When C = 9 And tol= 0.01 The result is: 0.9451932512256469
 When C = 9 And tol= 0.1 The result is: 0.934688695897694
 When C = 10 And tol= 0.0001 The result is: 0.9459939868291347
 When C = 10 And tol= 0.001 The result is: 0.946986854851645
 When C = 10 And tol= 0.01 The result is: 0.9453948258981593
 When C = 10 And tol= 0.1 The result is: 0.9352855087681448
 (0.9471844770612193, 8, 0.0001)
 When C = 1 And tol= 0.0001 The result is: 0.9501431651468467
 When C = 1 And tol= 0.001 The result is: 0.9503419551555095
 When C = 1 And tol= 0.01 The result is: 0.9483542340860097
 When C = 1 And tol= 0.1 The result is: 0.9404285464331421
 When C = 2 And tol= 0.0001 The result is: 0.9533284958623487
 When C = 2 And tol= 0.001 The result is: 0.9537225708544715
 When C = 2 And tol= 0.01 The result is: 0.9509388024529437
 When C = 2 And tol= 0.1 The result is: 0.9443906170222032
 When C = 3 And tol= 0.0001 The result is: 0.9547209882113634
 When C = 3 And tol= 0.001 The result is: 0.9545158334101689
 When C = 3 And tol= 0.01 The result is: 0.9527249642965824
 When C = 3 And tol= 0.1 The result is: 0.94478905008569
 When C = 4 And tol= 0.0001 The result is: 0.9567122611676098
 When C = 4 And tol= 0.001 The result is: 0.9561138469069848
 When C = 4 And tol= 0.01 The result is: 0.9537229574550018
 When C = 4 And tol= 0.1 The result is: 0.9463779420990039
 When C = 5 And tol= 0.0001 The result is: 0.956112284602273
 When C = 5 And tol= 0.001 The result is: 0.9567055674963665
 When C = 5 And tol= 0.01 The result is: 0.9533249487011378
 When C = 5 And tol= 0.1 The result is: 0.944993381385299
 When C = 6 And tol= 0.0001 The result is: 0.9565095044577431
 When C = 6 And tol= 0.001 The result is: 0.9569075476943864
 When C = 6 And tol= 0.01 The result is: 0.9527257579379279
 When C = 6 And tol= 0.1 The result is: 0.9471727824437071
 When C = 7 And tol= 0.0001 The result is: 0.9561099286373193
 When C = 7 And tol= 0.001 The result is: 0.956512284602273
 When C = 7 And tol= 0.01 The result is: 0.9539182502869427
 When C = 7 And tol= 0.1 The result is: 0.9461854245574903
 When C = 8 And tol= 0.0001 The result is: 0.9565122830400604
 When C = 8 And tol= 0.001 The result is: 0.9569055611912024

```

When C = 8 And tol= 0.01 The result is: 0.9543178510853458
When C = 8 And tol= 0.1 The result is: 0.9453834882344185
When C = 9 And tol= 0.0001 The result is: 0.9569087171026658
When C = 9 And tol= 0.001 The result is: 0.9567099286373193
When C = 9 And tol= 0.01 The result is: 0.9537230027713776
When C = 9 And tol= 0.1 The result is: 0.9463771577290423
When C = 10 And tol= 0.0001 The result is: 0.9569095357408866
When C = 10 And tol= 0.001 The result is: 0.9573071484927898
When C = 10 And tol= 0.01 The result is: 0.9543186510318558
When C = 10 And tol= 0.1 The result is: 0.9477680456562334
(0.9573071484927898, 10, 0.001)

```

So the result above for family class, the best classification rate is 93.7%. For Genus class, the classification rate is 94.7%, and for Species, the rate is 95.7%.

6 Repeat 1(b)iii by using SMOTE or any other method you know to remedy class imbalance. Report your conclusions about the classifiers you trained.

```
In [11]: from imblearn.over_sampling import SMOTE
```

```
In [26]: def getPerecentInList(alist):
        adict={}
        for number in alist:
            if number not in adict:
                adict[number]=1
            else:
                adict[number]+=1

        for key in adict:
            adict[key]=(adict[key])/len(alist)

        return adict
```

```
print(getPerecentInList(family_y))
```

```
{3: 0.6143154968728284, 1: 0.07533009034051424, 2: 0.30090340514246006, 0: 0.00945100764419736}
```

```
In [37]: #from the smote algorithm
        #family, genus, species
        #smote algoirithm will increase the imbalance data by
        sm = SMOTE(random_state=12)

        family_train_x_sm,family_train_y_sm=sm.fit_sample(family_x_train,family_y_train)
        genus_train_x_sm,genus_train_y_sm=sm.fit_sample(genus_x_train,genus_y_train)
        species_train_x_sm,species_train_y_sm=sm.fit_sample(species_x_train,species_y_train)
```

```
In [36]: #baseline is 92%
```

```
svm_model=svm.LinearSVC()  
svm_model.fit(family_x_train,family_y_train)  
family_test_pre_y=svm_model.predict(family_x_test)  
print("Accuracy Score is: "+str(accuracy_score(family_test_pre_y,family_y_test)))
```

Accuracy Score is: 0.9277443260768874

```
In [38]: svm_model=svm.LinearSVC()  
svm_model.fit(family_train_x_sm,family_train_y_sm)  
family_test_pre_y=svm_model.predict(family_x_test)  
print("Accuracy Score is: "+str(accuracy_score(family_test_pre_y,family_y_test)))
```

Accuracy Score is: 0.8985641500694767

```
In [39]: svm_model=svm.LinearSVC()  
svm_model.fit(genus_train_x_sm,genus_train_y_sm)  
score=svm_model.score(genus_x_test,genus_y_test)  
print(score)
```

0.9018063918480778

```
In [40]: svm_model=svm.LinearSVC()  
svm_model.fit(species_train_x_sm,species_train_y_sm)  
score=svm_model.score(species_x_test,species_y_test)  
print(score)
```

0.9620194534506716

From the above experiments, for some classification situation, SMOTE algorithm improves the accuracy. But some classes like family, it did not behave very well.

7 v. Extra Practice: Study the Classifier Chain method and apply it to the above problem.

```
In [32]: from sklearn.multioutput import ClassifierChain  
from sklearn.model_selection import train_test_split  
from sklearn.multiclass import OneVsRestClassifier  
from sklearn.metrics import jaccard_similarity_score  
from sklearn.linear_model import LogisticRegression  
from sklearn.datasets import fetch_mldata
```

```
In [33]: #for the family situation  
ovr = OneVsRestClassifier(LogisticRegression())  
ovr.fit(family_x_train,family_y_train)
```



```

Out [33]: OneVsRestClassifier(estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False),
    n_jobs=1)

In [34]: Y_pred_ovr = ovr.predict(family_x_test)
    ovr_jaccard_score = jaccard_similarity_score(family_y_test, Y_pred_ovr)

In [35]: chains = [ClassifierChain(LogisticRegression(), order='random', random_state=i, cv=5)
    for i in range(10)]

In [38]: Y = label_binarize(Y, classes=[0, 1, 2, 3])

In [43]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=.3)

In [44]: chains = [ClassifierChain(LogisticRegression(), order='random', random_state=i)
    for i in range(10)]
    for chain in chains:
        chain.fit(X_train, Y_train)

    Y_pred_chains = np.array([chain.predict(X_test) for chain in
        chains])
    chain_jaccard_scores = [jaccard_similarity_score(Y_test, Y_pred_chain >= .5)
        for Y_pred_chain in Y_pred_chains]

    Y_pred_ensemble = Y_pred_chains.mean(axis=0)
    ensemble_jaccard_score = jaccard_similarity_score(Y_test,
        Y_pred_ensemble >= .5)

    model_scores = [ovr_jaccard_score] + chain_jaccard_scores
    model_scores.append(ensemble_jaccard_score)

In [45]: model_names = ('Independent',
    'Chain 1',
    'Chain 2',
    'Chain 3',
    'Chain 4',
    'Chain 5',
    'Chain 6',
    'Chain 7',
    'Chain 8',
    'Chain 9',
    'Chain 10',
    'Ensemble')


x_pos = np.arange(len(model_names))

# Plot the Jaccard similarity scores for the independent model, each of the

```

chains, and the ensemble (note that the vertical axis on this plot does # not begin at 0).

```
fig, ax = plt.subplots(figsize=(7, 4))
ax.grid(True)
ax.set_title('Classifier Chain Ensemble Performance Comparison')
ax.set_xticks(x_pos)
ax.set_xticklabels(model_names, rotation='vertical')
ax.set_ylabel('Jaccard Similarity Score')
ax.set_ylim([min(model_scores) * .9, max(model_scores) * 1.1])
colors = ['r'] + ['b'] * len(chain_jaccard_scores) + ['g']
ax.bar(x_pos, model_scores, alpha=0.5, color=colors)
plt.tight_layout()
plt.show()
```



From the above result we can find that the ensemble model has almost 92.5% accuracy

8 vi. Extra Practice: Research how confusion matrices, precision, recall, ROC, and AUC are defined for multi-label classification and compute them for the classifiers you trained in above.

Precision-recall curves are typically used in binary classification to study the output of a classifier. In order to extend the precision-recall curve and average precision to multi-class or multi-label classification, it is necessary to binarize the output. So we will calculate the average precision and recall, result

9 For family

```
In [25]: X=np.array(family_data.drop("Family",axis=1))
        Y=family_data["Family"]
```

```
In [26]: from sklearn.preprocessing import label_binarize
        Y = label_binarize(Y, classes=[0, 1, 2,3])
        n_classes = Y.shape[1]
```

```
In [27]: Y
```

```
Out[27]: array([[0, 0, 0, 1],
               [0, 0, 0, 1],
               [0, 0, 0, 1],
               ...,
               [0, 0, 1, 0],
               [0, 0, 1, 0],
               [0, 0, 1, 0]])
```

```
In [28]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=.3)
```

```
# We use OneVsRestClassifier for multi-label prediction
from sklearn.multiclass import OneVsRestClassifier
```

```
# Run classifier
classifier = OneVsRestClassifier(svm.LinearSVC(random_state=42))
classifier.fit(X_train, Y_train)
y_score = classifier.decision_function(X_test)
```

```
In [29]: from sklearn.metrics import precision_recall_curve
        from sklearn.metrics import average_precision_score
```

```
# For each class
precision = dict()
recall = dict()
average_precision = dict()
for i in range(n_classes):
    precision[i], recall[i], _ = precision_recall_curve(Y_test[:, i],
                                                         y_score[:, i])
    average_precision[i] = average_precision_score(Y_test[:, i], y_score[:, i])
```

```
# A "micro-average": quantifying score on all classes jointly
precision["micro"], recall["micro"], _ = precision_recall_curve(Y_test.ravel(),
                                                                y_score.ravel())
average_precision["micro"] = average_precision_score(Y_test, y_score,
                                                         average="micro")
print('Average precision score, micro-averaged over all classes: {0:0.2f}'
      .format(average_precision["micro"]))
```

Average precision score, micro-averaged over all classes: 0.96

```
In [30]: #then make the confusion matrix, it is same as the binary situation
X=family_data.drop("Family",axis=1)
Y=family_data["Family"]
svm_linear=svm.LinearSVC(random_state=41)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=.3)
svm_linear.fit(X_train,Y_train)
Y_test_pre=svm_linear.predict(X_test)
```

```
In [31]: from sklearn.metrics import confusion_matrix
print(confusion_matrix(Y_test_pre,Y_test))
```

```
[[ 0  0  0  0]
 [ 0 154 14 12]
 [ 13 12 597 23]
 [ 0  3  57 1274]]
```

10 Make ROC and AUC curve

```
In [18]: #then make the ROC and AUC Curve for the family class.
import numpy as np
import matplotlib.pyplot as plt
from itertools import cycle

from sklearn import svm, datasets
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier
from scipy import interp

# Import some data to play with
X = family_data.drop("Family",axis=1)
y = list(family_data["Family"])

# Binarize the output
y = label_binarize(y, classes=[0,1,2,3])
n_classes = y.shape[1]
```

```
In [19]: # Add noisy features to make the problem harder
random_state = np.random.RandomState(0)
n_samples, n_features = X.shape
X = np.c_[X, random_state.randn(n_samples, 2 * n_features)]
```

```
In [20]: # shuffle and split training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.4,
```

```

random_state=0)

# Learn to predict each class against the other
classifier = OneVsRestClassifier(svm.SVC(kernel='linear', probability=True,
                                         random_state=random_state))
y_score = classifier.fit(X_train, y_train).decision_function(X_test)

In [21]: # Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Compute micro-average ROC curve and ROC area
fpr["micro"], tpr["micro"], _ = roc_curve(y_test.ravel(), y_score.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

In [22]: y_score = classifier.fit(X_train, y_train).decision_function(X_test)

# Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Compute micro-average ROC curve and ROC area
fpr["micro"], tpr["micro"], _ = roc_curve(y_test.ravel(), y_score.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

In [23]: all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))

# Then interpolate all ROC curves at this points
mean_tpr = np.zeros_like(all_fpr)
for i in range(n_classes):
    mean_tpr += interp(all_fpr, fpr[i], tpr[i])

# Finally average it and compute AUC
mean_tpr /= n_classes

fpr["macro"] = all_fpr
tpr["macro"] = mean_tpr
roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

```

```

# Plot all ROC curves
plt.figure()
plt.plot(fpr["micro"], tpr["micro"],
         label='micro-average ROC curve (area = {0:0.2f})'
         ''.format(roc_auc["micro"]),
         color='deeppink', linestyle=':', linewidth=4)

plt.plot(fpr["macro"], tpr["macro"],
         label='macro-average ROC curve (area = {0:0.2f})'
         ''.format(roc_auc["macro"]),
         color='navy', linestyle=':', linewidth=4)

colors = cycle(['aqua', 'darkorange', 'cornflowerblue', "blue", "yellow", "green"])
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, #lw=lw,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Some extension of Receiver operating characteristic to multi-class')
plt.legend(loc="lower right")
plt.show()

```

